

# "New Year, New You" Portfolio Challenge

Jan. 12, 2026

---

What follows is a **high-level build plan**, ordered deliberately so that each step constrains the next and prevents architectural drift.

This is not a task list — it's an **assembly sequence**.

## Phase 0 — Fix the constraints (before writing code)

**Outcome:** Clear boundaries that prevent overbuilding.

We decide up front:

- one container
- one public service
- static-first
- one AI integration
- no user state
- no database

This ensures the system remains inspectable, fast, and contest-appropriate.

Only after this do we touch a framework.

---

## Phase 1 — Curate the content (before choosing UI)

**Outcome:** Authoritative, frozen portfolio data.

We:

1. Select a **small set of projects** (5–7 max)
2. Write a short, structured summary for each:
  - problem
  - approach
  - tradeoffs
  - repo link
3. Store this as plain data:
  - `projects.json` or `projects.yaml`
  - Markdown files for project deep dives

At this point:

- no AI
- no frontend
- no deployment

This becomes the *single source of truth*.

---

## Phase 2 — Build the static site shell

**Outcome:** A complete portfolio **without** AI.

We:

1. Choose a static framework (Astro / Next static)

2. Implement routes:

- `/`
- `/portfolio`
- `/project/[slug]`
- `/how-i-work`
- `/about`

3. Render everything from the curated data

If the project stopped here, it would already be a valid portfolio.

This step proves:

- navigation
- layout
- readability
- performance

No dynamic behavior yet.

---

## Phase 3 — Introduce the AI boundary

**Outcome:** A clean seam between content and AI.

We:

1. Add a tiny backend API:

- one endpoint for “intro”

- one endpoint for “ask”
2. Write **fixed system prompts**
  3. Inject portfolio summaries as context
  4. Return a single response per request

Critically:

- no conversation memory
- no user personalization
- no streaming
- no retries

This phase is about *containment*, not capability.

---

## Phase 4 — Connect frontend to AI (minimally)

**Outcome:** AI appears as an assistant, not a feature takeover.

We:

- add a button on `/home` → calls intro endpoint
- add a small form on `/ask-the-portfolio` → calls ask endpoint
- display responses as plain text

No animations. No typing effects. No avatars.

If AI fails, the site still works.

---

## Phase 5 — Containerize and deploy

**Outcome:** Live, reproducible submission.

We:

1. Write a single Dockerfile

2. Build locally

3. Deploy to **Google Cloud Run**

4. Confirm:

- public access

- fast cold start

- predictable URL

This locks the runtime environment.

---

## Phase 6 — Embed and submit

**Outcome:** Contest-ready artifact.

We:

1. Embed the Cloud Run app in the DEV post

- Add required label:

```
--labels dev-tutorial=devnewyear2026
```

- 2.

3. Link GitHub repos clearly

4. State exactly what is original vs prior work

No last-minute features.

---

## Assembly order (compressed view)

1. **Content first** (projects + writing)
2. **Static site** (fully functional without AI)
3. **AI API** (small, bounded, optional)
4. **UI hooks** (buttons and forms only)
5. **Container + Cloud Run**
6. **DEV embed + submission**

If you reverse any of these, complexity creeps in.

---

## Why this order matters

- Content constrains AI
- Static pages constrain UX
- API constraints prevent “chatbot sprawl”
- Deployment last avoids wasted work

This is how you build something **simple that stays simple**.